

# Chapter 10

## Views

# Topics in this Chapter

- What are Views For?
- View Retrievals
- View Updates
- Snapshots
- SQL Facilities

# Views

- A view is created by applying an expression of relational algebra to a base relvar
- A view is a derived or virtual relvar
- Updates applied to the base relvar are immediately visible to the view
- Updates to values in the view “really” only update the base relvar by reiterating the original algebraic expression



# What are Views For?

- A view can make subsequent queries much simpler (analogous to a macro)
- Views allow the same data to be seen by different users in different ways at the same time
- Views provide automatic security for hidden data
- Views can promote logical data independence

# Logical Data Indendence

- Logical data independence means the immunity of users and user programs to changes in the logical structure of the database
- Growth aspect: New attributes or relvars can be added without affecting existing user programs
- Restructuring aspect: If a base relvar is split, the “old” structure can be emulated by a view; the “old” programs can now access the view

# Base Relvars and Views

- The Principle of Interchangeability (of Base and Derived Relvars):

There are no arbitrary and unnecessary distinctions between base and derived relvars

- The Principle of Database Relativity:

From a user's point of view, base relvars and views are logically equivalent



# View Retrievals

- A relational expression is a relation-valued function
- It takes in relation(s) as arguments, and returns a relation
- A view is created using a relational expression, in which the relation returned is the view

# View Updates

- A view is a relvar, and so is updateable by definition
- Earlier discussions of view updates tended to limit them
- It is in principle not necessary to limit them, apart from constraint enforcement
- View updates must adhere to the Golden Rule: no relvar must ever be allowed to violate its own predicate



# A View Updating Mechanism

- View updates must apply to appropriate underlying persistent relvars
- INSERT and DELETE should be inverses of each other
- UPDATE can be treated as a DELETE and an INSERT, so long as no intermediate actions or checks permitted
- View updates should be compatible with all relational operators

# A View Updating Mechanism – UNION

- Let view UV be defined as follows:
- VAR UV VIEW  
    ( S WHERE STATUS > 25 ) UNION  
    ( S WHERE CITY = 'Paris' ) ;
- Insert the following tuple into the view:
- ( S6, Smith, 50, Rome)
- This will insert the tuple into the suppliers base relvar, and thus appear in the view
- Similar effect for DELETE and UPDATE

# A View Updating Mechanism – INTERSECT

- If both sides of the intersect evaluate to true, the operation applies to the base relvars successfully
- Analogous to the effect in the case of a Union, except that both conditions must be satisfied
- For example, the status must be  $> 25$ , and the city must be Paris



# A View Updating Mechanism – Difference

- A MINUS B
- The predicate is  $(PA) \text{ AND NOT } (PB)$ ;
- The new tuple must satisfy PA, and not PB, and if so, it is inserted into A, or deleted from it, depending on which is invoked
- On an update the result must satisfy PA and not PB, and then the old version is deleted from A, and the updated version is inserted into A

# A View Updating Mechanism – Restrict

- A where p, with a predicate of PA
- The predicate for V is (PA) AND p
- If the new tuple is presented to the view based on the RESTRICT, and it satisfies PA and p, it is inserted, deleted, or updated, as appropriate

# A View Updating Mechanism – Project

- Let the attributes of A, with predicate  $PA$ , be partitioned into two disjoint groups, X and Y
- Consider X and Y as single composite attributes
- For INSERT, Y must have an acceptable default value to be used when X is satisfied (or *vice versa*)
- For DELETE, all tuples satisfied by X are deleted, regardless of the value of Y
- For UPDATE, prior Y values are retained



# A View Updating Mechanism – Join

- To update a tuple of a view based on a join, it may be necessary to affect other tuples to preserve the view's predicate
- Earlier writers concluded that it was impermissible to update such views
- Mr. Date believes a better approach is to follow through the ancillary implications, rather than to avoid such updates
- Joins can be one to one, one to many, or many to many

# A View Updating Mechanism – Join

- Let  $J = A \text{ JOIN } B$ , where  $A$ ,  $B$ , and  $J$  have headings  $\{X,Y\}$ ,  $\{Y,Z\}$  AND  $\{X,Y,Z\}$
- The predicates are  $PA$ ,  $PB$ ,  $PJ$ , where  $PJ$  for  $J$  is  $PA(a) \text{ AND } PB(b)$ , where  $a$  is the  $A$  portion and  $b$  is the  $B$  portion
- For INSERT, the new tuple  $j$  must satisfy  $PJ$ . If the  $A$  portion of  $j$  does not appear in  $A$ , it is inserted, and likewise for the  $B$  portion in  $B$
- Similar for DELETE and UPDATE

# A View Updating Mechanism – Join – One to One

- A one to one join is more accurately called (zero-or-one) to (zero-or-one)
- For example, a separate supplier relvar whose only other attribute is favorite restaurant, and not all suppliers have one
- The view that joins supplier to suprest can permit inserts, deletes and updates that affect the underlying relvars in appropriate ways



# A View Updating Mechanism – Join – One to Many

- A one to many join is more accurately called (zero-or-one) to (zero-or-more)
- Operations on the view will affect the underlying relvars appropriately
- A tuple inserted into such a view will insert a tuple into the underlying relvar on the one side, as well as a tuple into the underlying relvar on the many side
- DELETE will occur based on constraints on underlying relvars

# A View Updating Mechanism – Join – Many to Many

- A many to many join is more accurately called (zero-or-more) to (zero-or-more)
- Operations on the view will affect the underlying relvars appropriately
- A tuple inserted into such a view will insert a tuple into the underlying relvars
- DELETE will occur based on constraints on underlying relvars
- Base relvars and views retain their predicates

# Snapshots

- Snapshots are derived relvars, but they are not views
- Snapshots are the result of a relational operation that generates a read only relvar that is refreshed periodically
- An important case of controlled redundancy, used often for point in time reporting
- Snapshots have come to be misnamed “materialized views,” an oxymoron



# SQL Facilities

- CREATE VIEW <view name> AS  
    <table exp>  
    [WITH [<qualifier>] CHECK OPTION ] ;
- WITH CHECK OPTION is used for INSERTs and SELECTs: will make them fail if integrity constraints are violated; the default is that they will succeed regardless
- Can only be used with updateable views, as defined in SQL, which is narrower

# SQL Facilities – Updateable Views in SQL

- The standard is more complex than any implementation of it; simplified, we say:
- Views defined as a restriction and/or projection of a single base table
- Views defined as a one to one join
- The many side of a one to many join
- UNION ALL or INTERSECT
- Other restrictions may apply